

# **Deep Learning for Hate Speech Detection in Tweets**

by

pinkesh.badjatiya , Shashank Gupta, Manish Gupta, Vasudeva Varma

in

*26th International World Wide Web Conference, 2017*  
(WWW-2017)

Perth, Western Australia

Report No: IIIT/TR/2017/-1



Centre for Language Technologies Research Centre  
International Institute of Information Technology  
Hyderabad - 500 032, INDIA  
April 2017



**Table 1:** Comparison of Various Methods (Embedding Size=200 for GloVe as well as for Random Embedding)

	Method	Prec	Recall	F1
<b>Part A:</b> Baselines	Char n-gram+Logistic Regression [6]	0.729	0.778	0.753
	TF-IDF+Balanced SVM	0.816	0.816	0.816
	TF-IDF+GBDT	0.819	0.807	0.813
	BoWV+Balanced SVM	0.791	0.788	0.789
	BoWV+GBDT	0.800	0.802	0.801
<b>Part B:</b> DNNs Only	CNN+Random Embedding	0.813	0.816	0.814
	CNN+GloVe	0.839	0.840	0.839
	FastText+Random Embedding	0.824	0.827	0.825
	FastText+GloVe	0.828	0.831	0.829
	LSTM+Random Embedding	0.805	0.804	0.804
<b>Part C:</b> DNNs + GBDT Classi- er	LSTM+GloVe	0.807	0.809	0.808
	CNN+GloVe+GBDT	0.864	0.864	0.864
	CNN+Random Embedding+GBDT	0.864	0.864	0.864
	FastText+GloVe+GBDT	0.853	0.854	0.853
	FastText+Random Embedding+GBDT	0.886	0.887	0.886
	LSTM+GloVe+GBDT	0.849	0.848	0.848
	LSTM+Random Embedding+GBDT	<b>0.930</b>	<b>0.930</b>	<b>0.930</b>

(3) FastText: FastText [2] represents a document by average of word vectors similar to the BoWV model, but allows update of word vectors through Back-propagation during training as opposed to the static word representation in the BoWV model, allowing the model to fine-tune the word representations according to the task.

All of these networks are trained (fine-tuned) using labeled data with back-propagation. Once the network is learned, a new tweet is tested against the network which classifies it as racist, sexist or neither. Besides learning the network weights, these methods also learn task-specific word embeddings tuned towards the hate speech labels. Therefore, for each of the networks, we also experiment by using these embeddings as features and various other classifiers like SVMs and GBDTs as the learning method.

### 3. EXPERIMENTS

#### 3.1 Dataset and Experimental Settings

We experimented with a dataset of 16K annotated tweets made available by the authors of [6]. Of the 16K tweets, 3383 are labeled as sexist, 1972 as racist, and the remaining are marked as neither sexist nor racist. For the embedding based methods, we used the GloVe [5] pre-trained word embeddings. GloVe embeddings<sup>2</sup> have been trained on a large tweet corpus (2B tweets, 27B tokens, 1.2M vocab, uncased). We experimented with multiple word embedding sizes for our task. We observed similar results with different sizes, and hence due to lack of space we report results using embedding size=200. We performed 10-Fold Cross Validation and calculated weighted macro precision, recall and F1-scores.

We use 'adam' for CNN and LSTM, and 'RMS-Prop' for FastText as our optimizer. We perform training in batches of size 128 for CNN & LSTM and 64 for FastText. More details on the experimental setup can be found from our publicly available source code<sup>3</sup>.

#### 3.2 Results and Analysis

Table 1 shows the results of various methods on the hate speech detection task. Part A shows results for baseline methods. Parts B and C focus on the proposed methods where part B contains methods using neural networks only, while part C uses average of word embeddings learned by DNNs as features for GBDTs. We experimented with mul-

<sup>2</sup><http://nlp.stanford.edu/projects/glove/>

<sup>3</sup><https://github.com/pinkeshbadjatiya/twitter-hatespeech>

**Table 2:** Embeddings learned using DNNs clearly show the 'racist' or 'sexist' bias for various words.

Target Word	Similar words using GloVe	Similar words using task-specific embeddings learned using DNNs
pakistan	karachi, pakistani, lahore, india, taliban, punjab, islamabad	mohammed, murderer, pedophile, religion, terrorism, islamic, muslim
female	male, woman, females, women, girl, other, artist, girls, only, person	sexist, feminists, feminism, bitch, feminist, blonde, bitches, dumb, equality, models, cunt
muslims	christians, muslim, hindus, jews, terrorists, islam, sikhs, extremists, non-muslims, buddhists	islam, prophet, quran, slave, jews, slavery, pedophile, terrorist, terrorism, hamas, murder

iple classifiers but report results mostly for GBDTs only, due to lack of space.

As the table shows, our proposed methods in part B are significantly better than the baseline methods in part A. Among the baseline methods, the word TF-IDF method is better than the character n-gram method. Among part B methods, CNN performed better than LSTM which was better than FastText. Surprisingly, initialization with random embeddings is slightly better than initialization with GloVe embeddings when used along with GBDT. Finally, part C methods are better than part B methods. The best method is 'LSTM + Random Embedding + GBDT' where tweet embeddings were initialized to random vectors, LSTM was trained using back-propagation, and then learned embeddings were used to train a GBDT classifier. Combinations of CNN, LSTM, FastText embeddings as features for GBDTs did not lead to better results. Also note that the standard deviation for all these methods varies from 0.01 to 0.025.

To verify the task-specific nature of the embeddings, we show top few similar words for a few chosen words in Table 2 using the original GloVe embeddings and also embeddings learned using DNNs. The similar words obtained using deep neural network learned embeddings clearly show the 'hatred' towards the target words, which is in general not visible at all in similar words obtained using GloVe.

### 4. CONCLUSIONS

In this paper, we investigated the application of deep neural network architectures for the task of hate speech detection. We found them to significantly outperform the existing methods. Embeddings learned from deep neural network models when combined with gradient boosted decision trees led to best accuracy values. In the future, we plan to explore the importance of the user network features for the task.

### 5. REFERENCES

- [1] N. Djuric, J. Zhou, R. Morris, M. Grbovic, V. Radosavljevic, and N. Bhamidipati. Hate Speech Detection with Comment Embeddings. In *WWW*, pages 29{30, 2015.
- [2] A. Joulin, E. Grave, P. Bojanowski, and T. Mikolov. Bag of Tricks for Efficient Text Classification. *arXiv preprint arXiv:1607.01759*, 2016.
- [3] Y. Kim. Convolutional Neural Networks for Sentence Classification. In *EMNLP*, pages 1746{1751, 2014.
- [4] C. Nobata, J. Tetreault, A. Thomas, Y. Mehdad, and Y. Chang. Abusive Language Detection in Online User Content. In *WWW*, pages 145{153, 2016.
- [5] J. Pennington, R. Socher, and C. D. Manning. GloVe: Global Vectors for Word Representation. In *EMNLP*, volume 14, pages 1532{43, 2014.
- [6] Z. Waseem and D. Hovy. Hateful Symbols or Hateful People? Predictive Features for Hate Speech Detection on Twitter. In *NAACL-HLT*, pages 88{93, 2016.